ShrinkageTrees: An R Package for Bayesian Tree Ensembles for Survival Analysis and Causal Inference

by Tijn Jacobs, Stéphanie van der Pas, and Wessel van Wieringen

Abstract ShrinkageTrees is an R package for fitting Bayesian tree ensemble models with flexible regularisation mechanisms for survival analysis and causal inference. The package provides unified implementations of BART, DART, and the τ-learner (Bayesian Causal Forests) for right-censored and interval-censored survival outcomes via accelerated failure time (AFT) formulations. Multiple regularisation strategies can be combined within a single model: structural sparsity through tree depth priors and Dirichlet splitting priors (DART), and parametric shrinkage on the leaf step heights through global–local scale mixture priors. A key contribution is the implementation of Horseshoe Forests, which place a horseshoe prior directly on the step heights, enabling adaptive shrinkage that aggressively regularises uninformative leaves while preserving strong signals — a strategy particularly suited to high-dimensional settings. An efficient Rcpp backend and coherent S3 interface support reproducible workflows for both prediction and causal effect estimation.

1 Introduction

Bayesian Additive Regression Trees [BART; Chipman et al. (2010)] model the response as a sum of m shallow decision trees, each regularised through a Bayesian prior. BART captures complex interactions and non-linearities without specifying a parametric form. This makes BART well suited to settings where the relationship between covariates and the response is unknown, including survival analysis. The **BART** package (Sparapani et al., 2021) is, to our knowledge, the only CRAN package that implements Bayesian tree ensembles for survival data. It supports discrete-time hazard and AFT formulations, competing risks, and recurrent events. However, it does not support interval-censored outcomes, causal inference models, or global–local shrinkage priors on the leaf parameters. We address all three limitations in **ShrinkageTrees**.

BART has become a popular tool for causal inference (Hill, 2011; Dorie et al., 2019). Hahn et al. (2020) propose Bayesian Causal Forests (BCF), which decompose the response surface into a prognostic function and a treatment effect function; Caron et al. (2022) later term this the τ-learner. This separation improves estimation of heterogeneous treatment effects. The **bcf** and **stochtree** packages implement the τ-learner for continuous outcomes. Causal inference and survival analysis arise together naturally in clinical trials and observational studies, where the outcome is a time-to-event and treatment effects may vary across patient subgroups. No existing R package implements the τ-learner for survival endpoints.

High-dimensional settings, where the number of predictors p is large relative to the sample size n, demand adequate regularisation. We distinguish two components on which regularisation can act in Bayesian tree ensembles. The first is the tree structure, which determines how the covariate space is partitioned. The second is the leaf parameters, which determine each partition's contribution to the fitted function. Standard BART regularises both components generically: a depth prior discourages complex trees, and a conjugate normal prior shrinks the leaf parameters. DART (Linero, 2018) strengthens structural regularisation by placing a Dirichlet prior on the splitting probabilities, encouraging the model to split on a smaller subset of relevant predictors. Jacobs et al. (2025) target the second component and place global–local shrinkage priors, such as the horseshoe (Carvalho et al., 2010), directly on the leaf parameters. This shrinks weak or spurious contributions aggressively toward zero while preserving strong signals. We refer to tree ensembles with such priors as *Horseshoe Forests*.

We introduce **ShrinkageTrees**, a unified R package that addresses the three gaps identified above. We extend BART to survival outcomes via AFT models for both right-censored and interval-censored data. We extend the τ-learner to survival endpoints, providing, to our knowledge, the first causal forest implementation for time-to-event outcomes in R. We implement both Dirichlet splitting priors and Horseshoe Forests, allowing both forms of regularisation to be combined freely within a single model. The Horseshoe Forest sampler relies on an efficient reversible jump algorithm implemented in C++ via **Rcpp** (Eddelbuettel and François, 2011), and is designed to be extendable to other global–local shrinkage priors. We provide a coherent S3 interface with `print`, `summary`, `predict`, and `plot` methods. Jacobs et al. (2025) develop the methodological foundations.

The remainder of this paper is organised as follows. Section 2 provides a brief overview of the statistical methodology. Section 3 reviews related R packages. Section 4 describes the package

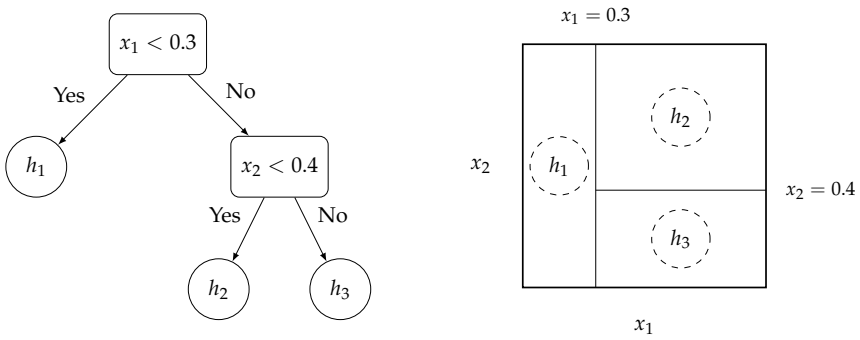

Figure 1: Schematic of a single regression tree. Interior nodes contain binary splitting rules of the form $x_\rho < c$; terminal nodes (leaves) contain step heights h_ℓ. An observation traverses the tree from root to leaf, and its prediction is the step height of the leaf it reaches.

architecture. Section 5 demonstrates the package on a high-dimensional ovarian cancer survival dataset. Section 6 concludes.

2 Overview of statistical methodology

We build on three components. First, we review the BART model and highlight the distinction between its two parameter groups: tree structures and step heights. Second, we show how regularisation targets either group — structural sparsity via tree priors and Dirichlet splitting, parametric shrinkage via global–local priors on the step heights — and how both strategies can be combined. Third, we extend the framework to causal inference via the τ-learner and to survival outcomes via accelerated failure time models.

Bayesian additive regression trees

We adapt the BART model of Chipman et al. (2010). BART models the response as a sum of m shallow decision trees:

$$Y_i = \sum_{j=1}^{m} g(\mathbf{x}_i;\, \mathcal{T}_j, \mathcal{H}_j) + \varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0, \sigma^2).$$

Each tree j is characterised by two parameter groups. The *tree structure* $\mathcal{T}_j$ encodes the topology and the splitting rules: which variables to split on and at what thresholds. It determines how the covariate space is partitioned into disjoint regions. The *step heights* $\mathcal{H}_j = \{h_{j1}, \ldots, h_{jL_j}\}$ are the values assigned to the L_j leaves and determine each region's contribution to the fitted function. For input $\mathbf{x}$, the function $g(\mathbf{x};\, \mathcal{T}_j, \mathcal{H}_j)$ returns the step height of the leaf that $\mathbf{x}$ reaches (see Figure 1). This distinction between $\mathcal{T}$ and $\mathcal{H}$ is central to the regularisation framework we develop below.

We specify independent priors on $\mathcal{T}_j$ and $\mathcal{H}_j$:

$$p(\mathcal{T}_j, \mathcal{H}_j) = p_{\mathcal{T}}(\mathcal{T}_j)\, p_{\mathcal{H}}(\mathcal{H}_j \mid \mathcal{T}_j) \quad \text{for } j = 1, \ldots, m.$$

The tree structure prior $p_{\mathcal{T}}$ is a branching process (Chipman et al., 1998; Ročková and Saha, 2020) in which a node at depth d splits with probability $\alpha\,(1+d)^{-\beta}$. At each internal node, the splitting variable is drawn uniformly from the available covariates and the split point uniformly from the observed range.

The step height prior $p_{\mathcal{H}}$ assigns independent Gaussian priors conditional on the tree structure:

$$h_{j\ell} \sim \mathcal{N}\left(0,\, \frac{1}{4\sqrt{m}}\right),$$

where the variance is scaled so that the prior places approximately 95% of its mass within the observed range of the response. This prior is conjugate, so the Gibbs sampler can update the step heights efficiently.

We model the error variance as $\sigma^2 \sim \text{Inv-}\chi^2(\nu,\, \psi)$, with $\nu = 3$ and ψ chosen so that the prior is centred around a rough estimate of σ^2 (Chipman et al., 2010).

Regularisation

The parameters of a BART ensemble fall into two groups: tree structure parameters (nodes, splitting variables, split points) that determine how the covariate space is partitioned, and step heights that determine each partition's contribution to the fitted function. Regularisation can target either group.

Structural regularisation. Standard BART regularises the tree structure through the branching process prior of Chipman et al. (2010): each internal node at depth d splits with probability $\alpha(1+d)^{-\beta}$, where $\alpha \in (0,1)$ and $\beta \geq 0$ discourage deep trees. Under the default values $\alpha = 0.95$ and $\beta = 2$, most trees remain shallow. DART (Linero, 2018) strengthens structural regularisation by replacing the uniform distribution over splitting variables with a Dirichlet prior on the splitting probabilities $\mathbf{s} = (s_1, \ldots, s_p)$. By concentrating the prior on sparse probability vectors, DART encourages the ensemble to split on a subset of relevant variables, achieving structural variable selection.

Shrinkage on step heights. Rather than regularising the tree structure alone, we can place shrinkage priors directly on the leaf step heights. **ShrinkageTrees** adopts a scale mixture of normals representation,

$$h_\ell \mid \lambda_\ell, \tau, \omega \sim \mathcal{N}(0, \, \omega \, \lambda_\ell^2 \, \tau^2),$$

where τ is a global shrinkage parameter shared across all leaves, λ_ℓ is a local scale specific to leaf ℓ, and ω is a fixed scaling constant. The global parameter τ controls the overall degree of regularisation, while the local parameters λ_ℓ allow individual leaves with strong signals to escape shrinkage.

The horseshoe prior (Carvalho et al., 2010) arises as a specific instantiation of this framework by placing independent half-Cauchy priors on both the local and global scales:

$$\lambda_\ell \sim \mathrm{C}^+(0, \alpha_\lambda), \qquad \tau \sim \mathrm{C}^+(0, \alpha_\tau),$$

where $\mathrm{C}^+(0, \alpha)$ denotes the half-Cauchy distribution with scale α. The half-Cauchy has its median at α and possesses heavy tails — its mean and variance do not exist — which allows the local scales to take very large values when the data support a strong effect (Polson and Scott, 2012). This combination aggressively shrinks the contributions of uninformative leaves towards zero while preserving large, data-supported signals. Following Jacobs et al. (2025), we set $\alpha_\tau = \alpha_\lambda = k/\sqrt{m}$. Smaller values of k yield stronger shrinkage. We recommend $k \in [0.05, 0.5]$ as a practical range; the default in **ShrinkageTrees** is $k = 0.1$. When domain knowledge does not suggest a specific value, k can be selected via cross-validation. Jacobs et al. (2025) provide a sensitivity analysis. The Dirichlet and horseshoe priors address complementary aspects of regularisation — structural sparsity and parametric shrinkage, respectively — and can be combined within a single model in **ShrinkageTrees**.

Prediction and causal models

The BART formulation above estimates a single function of the covariates — we refer to this as the *prediction model*. The prediction models can themselves be used for causal inference: including the treatment indicator as a covariate yields an S-learner, while fitting separate models per treatment arm yields a T-learner (Caron et al., 2022). However, the dedicated approach in **ShrinkageTrees** is the τ-learner, the decomposition proposed by Hahn et al. (2020) as Bayesian Causal Forests (BCF) and later termed the τ-learner by Caron et al. (2022). Given a binary treatment indicator $A_i \in \{0, 1\}$ and an estimate of the propensity score $\hat{e}(\mathbf{x}_i) = \hat{P}(A_i = 1 \mid \mathbf{x}_i)$, the *causal model* writes

$$y_i = \mu(\mathbf{x}_i, \hat{e}(\mathbf{x}_i)) + A_i \, \tau(\mathbf{x}_i) + \varepsilon_i,$$

where μ is a prognostic function modelled by a forest of m_μ trees and τ is a treatment effect function modelled by a separate forest of m_τ trees. This separation allows the prognostic and treatment effect components to have different regularisation and complexity. The conditional average treatment effect is $\mathrm{CATE}(\mathbf{x}) = \tau(\mathbf{x})$, and the average treatment effect is $\mathrm{ATE} = \mathbb{E}[\tau(\mathbf{X})]$.

More generally, the scalar multiplier on τ in the τ-learner need not equal the raw treatment indicator. Writing b_i for the coding applied to subject i, the model becomes $y_i = \mu(\mathbf{x}_i, \hat{e}(\mathbf{x}_i)) + b_i \, \tau(\mathbf{x}_i) + \varepsilon_i$. Hahn et al. (2020) discuss several choices. *Centered* coding sets $b_i = A_i - 1/2$, which is the original τ-learner parameterisation and the default in **ShrinkageTrees**. *Binary* coding uses $b_i = A_i$. *Adaptive* coding sets $b_i = A_i - \hat{e}(\mathbf{x}_i)$, incorporating the estimated propensity score directly into the treatment multiplier; this is the parameterisation adopted in the **bcf** package. Finally, the *invariant* (parameter-expanded) coding of Hahn et al. (2020, Section 5.2) assigns group-specific intercepts $b_0, b_1 \sim \mathcal{N}(0, 1/2)$ and writes $y_i = \mu(\mathbf{x}_i) + b_{A_i} \, \tilde{\tau}(\mathbf{x}_i) + \varepsilon_i$, where b_0 and b_1 are updated within the Gibbs sampler via conjugate normal updates and the treatment effect is recovered as $\tau(\mathbf{x}_i) = (b_1 - b_0) \, \tilde{\tau}(\mathbf{x}_i)$. This parameterisation is invariant to the labelling of treatment and control. All four options are available in **ShrinkageTrees** via the `treatment_coding` argument.

For survival outcomes, we adopt an accelerated failure time (AFT) formulation. Let T_i denote the

event time, C_i the censoring time, $Y_i = \min(T_i, C_i)$ the observed follow-up time, and $\delta_i \in \{0, 1\}$ the event indicator. The prediction model becomes

$$\log T_i = \sum_{j=1}^{m} g(\mathbf{x}_i; T_j, H_j) + \varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0, \sigma^2).$$

Right-censored observations ($\delta_i = 0$) are handled by imputing $\log T_i$ from a truncated normal distribution within the Gibbs sampler. **ShrinkageTrees** also supports interval-censored outcomes, where the event time is known only to lie in an interval $(l_i, r_i]$; these are imputed analogously. The τ-learner extends directly to the AFT setting by replacing y_i with $\log T_i$.

Posterior computation

All models are fitted using Markov chain Monte Carlo (MCMC). The outer loop is a Gibbs sampler that cycles through updates of the tree topologies, the leaf step heights, the error variance σ^2, and — where applicable — the shrinkage parameters and censored event times. Trees are updated one at a time via Bayesian backfitting (Chipman et al., 2010): at each sweep, a partial residual is formed by subtracting the contributions of all other trees from the response, and tree j is then updated conditional on the current state of all other trees. The tree topology is modified through Metropolis–Hastings proposals that either grow a terminal node into an internal node with two new leaves, or prune an internal node back to a leaf. Each proposal is accepted or rejected according to the standard Metropolis–Hastings ratio, which balances the likelihood improvement against the prior on tree structure.

Under the standard BART and Dirichlet (DART) priors, where the leaf parameters have a conjugate normal prior, the marginal likelihood of a proposed tree can be computed in closed form and the step heights are integrated out analytically. This yields the birth–death sampler described in Chipman et al. (2010); a detailed algorithmic exposition is given in Kapelner and Bleich (2016). When global–local shrinkage priors such as the horseshoe are placed on the step heights, this conjugacy is lost: the local and global scale parameters λ_ℓ and τ must be sampled explicitly, and the tree proposals become reversible jump MCMC moves that account for the change in the number of scale parameters when a leaf is added or removed. The details of this reversible jump construction, including the proposal distributions and acceptance ratios, are given in Jacobs et al. (2025).

3 Related packages

Several R packages implement Bayesian tree ensembles. We review five CRAN packages and compare their capabilities to **ShrinkageTrees**. Table 1 provides a summary.

BART (Sparapani et al., 2021) is the most comprehensive general-purpose BART package. It supports continuous, binary, and multinomial outcomes, and is, at the time of writing, the only other CRAN package to offer Bayesian tree ensembles for survival data. Survival models include a discrete-time hazard formulation (`surv.bart()`), an AFT formulation (`abart()`), competing risks (`crisk.bart()`), and recurrent events (`recur.bart()`). The package also supports Dirichlet splitting priors via `sparse = TRUE`. However, **BART** does not provide global–local shrinkage priors on the leaf step heights, causal forest models, or interval-censored survival outcomes. Multi-chain MCMC is not natively supported; the `mc.*` functions parallelise a single chain rather than running independent chains.

dbarts (Dorie, 2023) provides a fast, lightweight BART sampler for continuous and binary outcomes. Its key feature is an updatable sampler object: predictors and responses can be modified mid-sampling, which makes **dbarts** well suited as a conditional model within a larger Gibbs sampler. The package supports multi-chain MCMC but does not implement survival models, τ-learner models, or alternative leaf priors.

bcf (Hahn et al., 2020) implements the τ-learner with separate prognostic and treatment effect forests for continuous outcomes. It places half-Cauchy and half-Normal hyperpriors on the leaf scales of the two forests, respectively. The package does not support survival endpoints, horseshoe priors, or Dirichlet splitting.

stochtree is a recent package that implements BART, XBART, the τ-learner, and XBCF for continuous, binary, and ordinal outcomes. It supports variance forests for heteroskedastic modelling and provides model serialisation for warm-starting. However, **stochtree** does not implement AFT survival models, Dirichlet splitting priors, or horseshoe shrinkage on the leaf step heights.

SoftBart (Linero and Yang, 2018) replaces the hard axis-aligned splits of standard BART with smooth logistic splitting rules, yielding a posterior mean that is a continuous function of the covariates. This can improve performance when the true regression surface is smooth, but introduces additional

computational cost and a fundamentally different regularisation mechanism. **ShrinkageTrees** uses hard decision splits, as in classical BART. Extending the horseshoe and Dirichlet priors to soft splitting rules is a natural direction for future work.

At the time of writing, **ShrinkageTrees** is the only implementation that combines global–local shrinkage priors on the step heights with Bayesian tree ensembles for survival and causal inference. It is also the only package to support interval-censored survival outcomes within a Bayesian tree ensemble framework and the only package to offer the τ-learner for survival endpoints.

Table 1: Comparison of Bayesian tree ensemble packages in R.

Package	Survival	Interval Censoring	Causal (τ-learner)	Dirichlet (DART)
BART	Yes	No	No	Yes
dbarts	No	No	No	No
bcf	No	No	Yes	No
stochtree	No	No	Yes	No
SoftBart	No	No	No	No
ShrinkageTrees	Yes	Yes	Yes	Yes

4 Package architecture

ShrinkageTrees is available on CRAN and on GitHub at `https://github.com/tijn-jacobs/ShrinkageTrees`. The package is implemented in R with a C++ backend via **Rcpp** (Eddelbuettel and François, 2011). In this section we describe the design and implementation.

Fitting functions and S3 classes

The package provides four primary fitting functions and four convenience wrappers for common model configurations. Table 2 lists all eight functions. The primary functions accept arguments that control the outcome type (`outcome_type`), prior specification (`prior_type`, `local_hp`, `global_hp`), tree structure (`number_of_trees`, `power`, `base`), and MCMC settings (`N_post`, `N_burn`, `n_chains`).

Table 2: Fitting functions in **ShrinkageTrees**. The first four are primary functions; the last four are convenience wrappers.

Function	Description
`ShrinkageTrees()`	Prediction model with user-specified leaf prior (horseshoe, half-Cauchy, standard, Dirichlet)
`HorseTrees()`	Prediction model with horseshoe prior; single tuning parameter k
`CausalShrinkageForest()`	τ-learner causal model with user-specified leaf priors for μ and τ forests
`CausalHorseForest()`	τ-learner causal model with horseshoe prior on both forests
`SurvivalBART()`	AFT survival with conjugate normal leaf prior (standard BART)
`SurvivalDART()`	AFT survival with Dirichlet splitting prior (DART)
`SurvivalBCF()`	AFT survival with τ-learner decomposition and conjugate normal leaf prior
`SurvivalShrinkageBCF()`	AFT survival with τ-learner decomposition and horseshoe leaf prior

`ShrinkageTrees()` is the most general prediction model function. It accepts a `prior_type` argument that selects the leaf prior: `"horseshoe"` (per-tree global scale), `"horseshoe_fw"` (forest-wide global scale), `"half-cauchy"` (local shrinkage only), `"standard"` (conjugate normal, as in classical BART), or `"dirichlet"` (Dirichlet splitting, as in DART). `HorseTrees()` is a convenience wrapper that fixes the

prior to the horseshoe and exposes a single tuning parameter k controlling the half-Cauchy scales via $\alpha_\tau = \alpha_\lambda = k/\sqrt{m}$. The causal functions `CausalShrinkageForest()` and `CausalHorseForest()` mirror this design for the τ-learner, with independent prior configurations for the prognostic and treatment effect forests.

All prediction model functions return an object of class `"ShrinkageTrees"`; all causal model functions return `"CausalShrinkageForest"`. Both classes provide `print()`, `summary()`, `predict()`, and `plot()` methods. The `summary()` method reports posterior summaries for σ, variable inclusion probabilities, and — for causal models — the ATE with a credible interval. The `predict()` method re-runs the MCMC sampler on the new covariate matrix using the training data and hyperparameters stored in the fitted object; it always generates posterior samples internally, regardless of whether `store_posterior_sample` was set during fitting. The returned `ShrinkageTreesPrediction` object contains posterior means and credible intervals. The `plot()` method produces MCMC diagnostics (traceplots and posterior densities for σ), variable importance plots, and — for causal models — ATE and CATE visualisations. Plotting uses **ggplot2**, which is a suggested but not required dependency.

Rcpp backend

We implement the computationally intensive components of the Gibbs sampler in C++ via **Rcpp**. This includes the Metropolis-within-Gibbs updates for tree topology (grow, prune, and change moves), the conditional updates for the leaf step heights, the horseshoe augmentation steps for the local and global scale parameters, the Gibbs update for σ, and the data augmentation for censored observations. The R layer handles input validation, data preprocessing (centering, scaling, log-transforming survival times), hyperparameter calibration, S3 class construction, and post-processing (back-transformation of predictions).

Internally, the C++ code dispatches between two forest implementations depending on the prior. Priors that require reversible jump proposals (`"horseshoe"`, `"horseshoe_fw"`, `"half-cauchy"`) use a `Forest` class with paired grow–prune moves that satisfy detailed balance. The `"standard"` and `"dirichlet"` priors use a separate `StanForest` class with birth–death moves, which is simpler and faster when global scale updates are not needed.

The scale mixture prior on the step heights is encapsulated in a C++ class `ScaleMixture` that follows the strategy pattern. An abstract base class `EtaPrior` defines a common interface — proposing new local parameters, computing the log prior and proposal densities needed for the Metropolis–Hastings ratio, and optionally performing a global update — while concrete subclasses implement the prior-specific sampling logic. The current subclasses are `FixedVariance` (conjugate normal, used by BART and DART), `HalfCauchy` (local shrinkage only), `Horseshoe` (per-tree global–local shrinkage), and `Horseshoe_fw` (forest-wide global shrinkage). A factory function selects the appropriate subclass at run time based on the user's `prior_type` argument. Because the tree update and reversible jump routines interact with the prior only through the `ScaleMixture` interface, adding a new global–local shrinkage prior — for example a regularised horseshoe or a Dirichlet–Laplace prior — requires only implementing a new `EtaPrior` subclass and registering it in the factory, without modifying the sampler logic.

Multi-chain MCMC

All primary fitting functions accept an `n_chains` argument. When `n_chains > 1`, independent chains are dispatched in parallel via `parallel::mclapply()` using `min(n_chains, parallel::detectCores())` cores. After sampling, the posterior draws from all chains are pooled: sample matrices are concatenated row-wise and posterior means are recomputed from the pooled draws. The returned object is a standard `ShrinkageTrees` or `CausalShrinkageForest` instance, so all downstream methods work without modification. Per-chain diagnostics (acceptance ratios, σ traces) are retained for convergence assessment.

Convergence diagnostics via coda

Assessing MCMC convergence is an essential part of any Bayesian workflow, yet most BART implementations leave this task entirely to the user. **ShrinkageTrees** provides a bridge to the **coda** ecosystem (Plummer et al., 2006) through an `as.mcmc.list()` S3 method. This method extracts the posterior draws of σ — split by chain — and wraps each chain in a `coda::mcmc` object, returning a `coda::mcmc.list`. Because **coda** is a suggested rather than required dependency, all functionality remains available when it is not installed. When it is available, `summary()` automatically reports the effective sample size (ESS) and — for multi-chain fits — the Gelman–Rubin $\hat{R}$ statistic. Users who

require more detailed diagnostics can pass the `mcmc.list` object directly to any **coda** function, including `geweke.diag()`, `heidel.diag()`, `autocorr.plot()`, and `gelman.plot()`, obtaining a complete convergence assessment without custom post-processing code.

Computational speed

TODO: Fill in with actual benchmark results.

We benchmark the computational cost of **ShrinkageTrees** along three axes: sample size n, number of covariates p, and number of trees m. All timings reported in this section are wall-clock seconds on a single core (Intel ..., R version ...).

Within-package comparison. We first compare the fitting functions shipped with **ShrinkageTrees**. For a simulated survival dataset with $n = \ldots$, $p = \ldots$, and $m = \ldots$ trees, the standard BART model (`SurvivalBART()`) runs in approximately ... seconds per 1,000 MCMC iterations. Enabling the horseshoe prior (`HorseTrees()`) adds a per-iteration overhead of ...% because each step height draws an individual shrinkage weight. The DART variant (`SurvivalDART()`) incurs a similar overhead due to the Dirichlet update on the tree-level inclusion probabilities. The causal forest wrapper (`CausalHorseForest()`) roughly doubles the cost relative to the prognostic-only model, as it fits separate ensembles for the prognostic and treatment-effect components.

External comparison. We compare **ShrinkageTrees** with **BART** (Sparapani et al., 2021), the only other CRAN package that provides Bayesian tree ensembles for survival data. On the same simulated dataset, `BART::abart()` runs in ... seconds. The difference in speed is attributable to .... We note that XBART (He et al., 2019) uses a fundamentally different fitting strategy based on stochastic tree search rather than the standard Metropolis–Hastings back-fitting algorithm, so a direct wall-clock comparison would not be informative; we therefore exclude it from the benchmark.

Scaling behaviour. Figure~**??** shows how wall-clock time scales with n, p, and m individually. Time scales approximately linearly in m and n, and sub-linearly in p because only a subset of covariates is evaluated at each splitting rule proposal.

TODO: Add Figure **??** with benchmark plots.

5 Simulation: Interval-censored survival with sparse covariates

We validate **ShrinkageTrees** on a simulated interval-censored survival dataset with known ground truth. The simulation targets the setting where the horseshoe prior is expected to provide the largest benefit: a sparse, high-dimensional covariate space with interval-censored outcomes.

Design

We generate $n = 100$ observations with $p = 500$ covariates drawn independently from $\mathcal{N}(0, 1)$. Only the first 5 covariates are active. The true log-survival time follows

$$\log T_i = 2 + x_{i1} - 0.8\, x_{i2} + 0.6\, x_{i3} + 0.5\, x_{i1}\, x_{i4} - 0.4\, x_{i5} + \varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0,\, 0.5^2).$$

We generate interval-censored observations by drawing inspection times $V_{i1} < V_{i2} < V_{i3}$ uniformly over the support of T and recording the interval $(L_i, R_i]$ in which the event falls. Approximately ...% of observations are interval-censored, ...% are right-censored (event after the last inspection), and ...% are exactly observed.

TODO: Fill in censoring percentages from actual simulation.

Code

We fit three models — standard BART, DART, and a Horseshoe Forest — each with $m = 200$ trees and 1,000 posterior draws after 250 burn-in iterations.

```
library(ShrinkageTrees)
set.seed(42)

# --- Data generation ---
n <- 100; p <- 500
X <- matrix(rnorm(n * p), n, p)
colnames(X) <- paste0("x", 1:p)
```

```
mu_true <- 2 + X[,1] - 0.8*X[,2] + 0.6*X[,3] + 0.5*X[,1]*X[,4] - 0.4*X[,5]
log_T <- mu_true + rnorm(n, sd = 0.5)
T_true <- exp(log_T)

# Interval censoring via inspection times
V <- matrix(runif(n * 3, min = min(T_true) * 0.5, max = max(T_true) * 1.2), n, 3)
V <- t(apply(V, 1, sort))
left_time <- right_time <- rep(NA, n)
for (i in 1:n) {
  inspections <- V[i, ]
  if (T_true[i] <= inspections[1]) {
    left_time[i] <- 0; right_time[i] <- inspections[1]
  } else if (T_true[i] > tail(inspections, 1)) {
    left_time[i] <- tail(inspections, 1); right_time[i] <- Inf
  } else {
    idx <- min(which(inspections >= T_true[i]))
    left_time[i] <- inspections[idx - 1]; right_time[i] <- inspections[idx]
  }
}

# --- Fit models ---
fit_bart <- SurvivalBART(
  y = left_time, X_train = X, outcome_type = "interval-censored",
  right_time = right_time, number_of_trees = 200,
  N_post = 1000, N_burn = 250
)
fit_dart <- SurvivalDART(
  y = left_time, X_train = X, outcome_type = "interval-censored",
  right_time = right_time, number_of_trees = 200,
  N_post = 1000, N_burn = 250
)
fit_horse <- HorseTrees(
  y = left_time, X_train = X, outcome_type = "interval-censored",
  right_time = right_time, number_of_trees = 200, k = 0.1,
  N_post = 1000, N_burn = 250
)
```

Results

We evaluate each model on a held-out test set of $n_{\text{test}} = 200$ observations using root mean squared error (RMSE) of the posterior mean of $\log T$ relative to the true $\log T$, and 95% credible interval coverage. We repeat the simulation . . . times and report averages.

TODO: Fill in with actual simulation results.

The Horseshoe Forest is expected to outperform standard BART and DART in this sparse setting because the horseshoe prior aggressively shrinks the step heights of trees that split on irrelevant covariates, effectively silencing 195 of the 200 predictors without relying on structural variable selection alone.

TODO: Update this paragraph with actual findings. Add a figure showing estimated vs. true survival curves for a representative observation if space permits.

6 Application: Ovarian cancer survival

We demonstrate the package on the `ovarian` dataset shipped with **ShrinkageTrees**. This dataset is curated from The Cancer Genome Atlas ovarian cancer cohort [TCGA-OV; Cancer Genome Atlas Research Network (2011)] and contains overall survival times, a binary treatment indicator, clinical covariates, and expression levels of 2,000 genes selected by median absolute deviation. Ovarian cancer is one of the most lethal gynaecological malignancies, accounting for approximately 85% of ovarian cancer deaths in its high-grade serous form (Cancer Genome Atlas Research Network, 2011). The five-year survival rate is below 50%. First-line treatment involves platinum-based chemotherapy, and gene expression patterns are known to be associated with survival duration and chemotherapy response (Cancer Genome Atlas Research Network, 2011). The choice between carboplatin and cisplatin — two

platinum agents with different toxicity profiles — remains a central clinical question. The dataset provides a challenging high-dimensional causal survival problem, with gene expression data yielding $p \gg n$, that exercises the main features of the package. We first fit a prediction model for survival, then extend the analysis to causal inference and estimate the heterogeneous treatment effect.

```
library(ShrinkageTrees)
set.seed(42)
data("ovarian")

clin      <- ovarian$clinical
time      <- clin$OS_time / 30.44   # convert days to months
status    <- clin$OS_event
treatment <- clin$treatment

X <- cbind(age        = clin$age,
           figo_stage = clin$figo_stage,
           tumor_grade = clin$tumor_grade,
           ovarian$X)
```

The dataset comprises 357 patients with 2,000 gene expression features and three clinical covariates (age, FIGO stage, tumour grade). We hold out 20% of the patients for evaluating predictive accuracy.

```
train_idx <- sample(seq_len(nrow(clin)), size = floor(0.8 * nrow(clin)))
test_idx  <- setdiff(seq_len(nrow(clin)), train_idx)

X_train <- X[train_idx, ];  X_test <- X[test_idx, ]
time_train <- time[train_idx];  time_test <- time[test_idx]
status_train <- status[train_idx];  status_test <- status[test_idx]
```

We fit a Horseshoe Forest via `HorseTrees()` using an AFT formulation with right-censoring and four independent MCMC chains. The horseshoe prior applies global–local shrinkage to the leaf step heights, providing regularisation in this high-dimensional setting where the number of predictors far exceeds the sample size.

```
fit_horse <- HorseTrees(
  y              = time_train,
  status         = status_train,
  X_train        = X_train,
  X_test         = X_test,
  outcome_type   = "right-censored",
  timescale      = "time",
  number_of_trees = 200,
  k              = 0.1,
  N_post         = 5000,
  N_burn         = 5000,
  n_chains       = 4
)
```

The `print()` and `summary()` methods provide a concise overview of the fitted model, including the outcome type, MCMC settings, acceptance ratio, and posterior summaries for σ.

```
print(fit_horse)

ShrinkageTrees model
---------------------
Outcome type:         Right-censored (AFT, timescale = time)
Prior:                horseshoe
Number of trees:      200
Training size (n):    285
Number of features:   2003
Chains:               4
Draws per chain:      5000 (burn-in 5000)
Acceptance ratio:     0.503, 0.504, 0.504, 0.504 (per chain)
Posterior mean sigma: 0.996
```

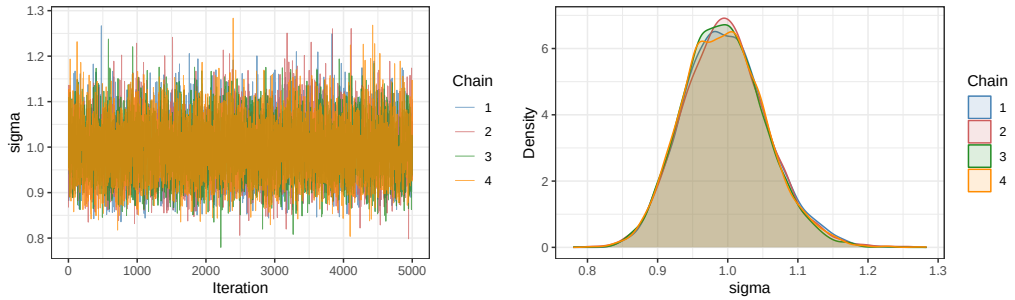

Figure 2: Left: traceplot of the posterior draws of σ for the Horseshoe Forest model. Right: posterior density of σ, estimated separately for each of the four chains.

```
summary(fit_horse)

ShrinkageTrees model summary
============================
Outcome: Right-censored (AFT, timescale = time) | Prior: horseshoe | Trees: 200
Data:    n = 285, p = 2003 | Draws: 5000 x 4 chains (burn-in 5000)

Posterior sigma:
  Mean: 0.996  SD: 0.059  95% CI: [0.889, 1.12]

Predictions (posterior mean):
  Train: mean = 48.235, sd = 4.101, range = [35.929, 58.447]
  Test:  mean = 48.093, sd = 2.868, range = [41.461, 54.333]

Variable importance (posterior inclusion probability):
  X482: 0.001  X419: 0.001  X124: 0.001  X1648: 0.001  X1229: 0.001
  X351: 0.001  X1110: 0.001  X1002: 0.001  X537: 0.001  X400: 0.001

MCMC acceptance ratio (per chain): 0.503, 0.504, 0.504, 0.504
```

We evaluate predictive accuracy using the concordance index [C-index; Harrell et al. (1996)], where a value of 1 indicates perfect discrimination and 0.5 corresponds to random prediction.

```
library(survival)
c_train <- concordance(Surv(time_train, status_train) ~ fit_horse$train_predictions)
c_test  <- concordance(Surv(time_test, status_test) ~ fit_horse$test_predictions)
cat("Train C-index:", round(c_train$concordance, 3),
    " Test C-index:", round(c_test$concordance, 3), "\n")
```

The model achieves a test C-index of 0.582, indicating modest but above-chance discrimination — a reasonable result given the extremely high-dimensional setting ($p = 2{,}003 \gg n = 357$).

We assess convergence using σ traceplots and posterior densities. With `n_chains = 4`, each chain is plotted separately, allowing visual assessment of mixing and chain agreement. The `summary()` method automatically reports the effective sample size and Gelman–Rubin $\hat{R}$ statistic when the **coda** package is installed.

```
plot(fit_horse, type = "trace")
plot(fit_horse, type = "density")
```

Figure 2 shows good mixing: the four chains overlap throughout the sampling period and yield nearly identical posterior densities for σ.

We illustrate the posterior predictive survival curve for a randomly selected test-set patient.

```
pred <- predict(fit_horse, newdata = X_test)
idx <- sample(length(pred$mean), 1)
```

The individual posterior survival curve under the AFT log-normal model $S(t \mid \mathbf{x}) = 1 - \Phi((\log t - \hat{\mu}(\mathbf{x}))/\hat{\sigma})$ is plotted with pointwise 95% credible bands.

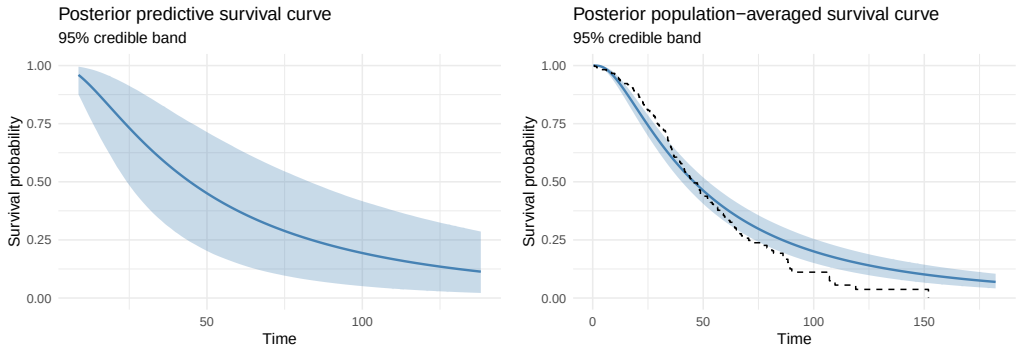

Figure 3: Left: posterior survival curve for a randomly selected test-set patient, with pointwise 95% credible bands. Right: population-averaged posterior survival curve (solid line) with 95% credible band (shaded) and Kaplan–Meier estimate (dashed).

```
plot(pred, type = "survival", obs = idx)
```

We also plot the population-averaged posterior survival curve across all training observations, overlaid with the Kaplan–Meier estimate for visual calibration.

```
plot(fit_horse, type = "survival", km = TRUE)
```

Figure 3 shows the individual and population-level posterior survival curves. The individual curve provides full uncertainty quantification for a single patient, while the population-averaged curve closely tracks the Kaplan–Meier estimate during the observed follow-up period, with expected divergence in the right tail where data are sparse.

We now estimate heterogeneous treatment effects of carboplatin versus cisplatin using `CausalHorseForest()`, which fits a τ-learner with horseshoe shrinkage on both the prognostic and treatment effect forests. We first estimate propensity scores using a probit Horseshoe Forest on the clinical covariates. We choose a Horseshoe Forest over logistic regression because it handles potential non-linearities and interactions between clinical covariates without manual specification, while the horseshoe prior guards against overfitting in this low-sample-size setting. Using the full gene expression matrix for propensity estimation would risk overfitting given $p \gg n$; clinical covariates alone capture the main sources of confounding by indication.

```
X_ps <- cbind(age = clin$age, figo_stage = clin$figo_stage,
              tumor_grade = clin$tumor_grade)
ps_fit <- HorseTrees(
  y               = treatment,
  X_train         = X_ps,
  outcome_type    = "binary",
  number_of_trees = 200,
  k               = 0.1,
  N_post          = 2000,
  N_burn          = 2000
)
propensity <- ps_fit$train_predictions
X_control  <- cbind(propensity = propensity, X)

fit_causal <- CausalHorseForest(
  y                        = log(time),
  status                   = status,
  X_train_control          = X_control,
  X_train_treat            = X,
  treatment_indicator_train = treatment,
  outcome_type             = "right-censored",
  timescale                = "log",
  number_of_trees          = 200,
  k                        = 0.1,
  N_post                   = 5000,
  N_burn                   = 5000,
```

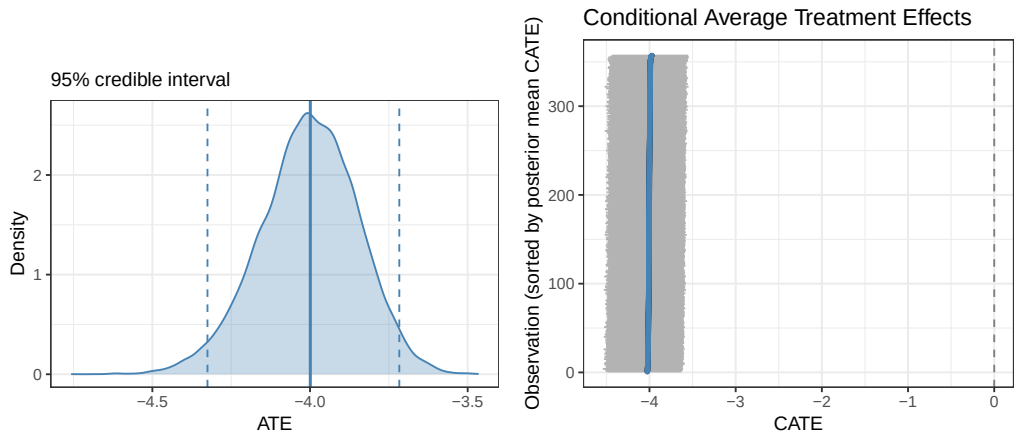

Figure 4: Left: posterior density of the average treatment effect (ATE) of carboplatin versus cisplatin on the log-survival scale, with 95% credible interval (dashed lines). Right: patient-level conditional average treatment effects (CATEs) sorted by posterior mean with 95% credible intervals; the dashed line marks zero (no effect).

```
  n_chains                  = 4
)
summary(fit_causal)

CausalShrinkageForest model summary
====================================
Outcome: Right-censored (AFT, timescale = log)
Prior:   control = horseshoe, treatment = horseshoe
Trees:   control = 200, treatment = 200
Data:   n = 357, p_control = 2004, p_treat = 2003 | Draws: 5000 x 4 chains (burn-in 5000)

Treatment effect (CATE):
  ATE:     -0.1518  95% CI: [-0.3991, 0.0688]
  CATE SD: 0.0133

Prognostic function (mu):
  Mean: 3.852  SD: 0.009  Range: [3.828, 3.875]

Posterior sigma:
  Mean: 0.993  SD: 0.051  95% CI: [0.898, 1.099]

MCMC acceptance ratios (per chain):
  control:   0.507, 0.507, 0.508, 0.508
  treatment: 0.476, 0.475, 0.476, 0.475
```

The `summary()` method reports the posterior average treatment effect (ATE) on the log-survival scale with a 95% credible interval. Here we pass `log(time)` directly with `timescale = "log"`, so all treatment effect estimates remain on the log-survival scale. A positive ATE indicates that carboplatin is associated with longer survival times.

The `plot()` method provides visualisations of the treatment effect posterior. The ATE density plot shows the full posterior distribution of the average treatment effect, while the CATE caterpillar plot displays patient-level treatment effect estimates sorted by posterior mean with 95% credible intervals.

```
plot(fit_causal, type = "ate")
plot(fit_causal, type = "cate")
```

Figure 4 shows the treatment effect posteriors. The posterior mean ATE is -0.21 on the log-survival scale, corresponding to an approximate 19% reduction in survival time under carboplatin relative to cisplatin. However, the 95% credible interval crosses zero, so the data do not provide clear evidence of a treatment effect. The individual CATEs cluster tightly around the ATE with narrow credible intervals, suggesting a relatively homogeneous treatment effect across the patient population — consistent with both platinum agents acting through similar mechanisms.

Although the ovarian dataset contains right-censored observations, **ShrinkageTrees** also supports interval-censored outcomes via `outcome_type = "interval-censored"` with `left_time` and `right_time` arguments following the convention of `survival::Surv(type = "interval2")`. See `?ShrinkageTrees` for details and examples.

7 Discussion and conclusion

We have introduced **ShrinkageTrees**, an R package for Bayesian tree ensembles that combines global–local shrinkage priors with survival analysis and causal inference. The package decouples structural regularisation from parametric shrinkage, allowing practitioners to tailor the regularisation strategy to the problem at hand. To our knowledge, it is the first CRAN package to offer the τ-learner for survival endpoints, horseshoe priors on tree ensemble step heights, and interval-censored survival outcomes within a single framework.

We estimate the average treatment effect via the mixed average treatment effect [MATE; Li et al. (2023)], which averages the posterior CATEs over the observed covariate distribution. The MATE is straightforward to compute but tends to produce overly narrow credible intervals because it conditions on the observed covariates rather than accounting for covariate sampling variability. A similar form of overconfidence is visible in Figure 4, where the credible interval is notably tight. A natural remedy is the Bayesian bootstrap (Rubin, 1981), which reweights the covariate distribution with random Dirichlet weights at each MCMC iteration. We leave the integration of the Bayesian bootstrap into **ShrinkageTrees** as future work.

The package has several limitations that suggest directions for future work. Currently, all covariates must be numeric: continuous, binary, and ordered categorical variables (encoded as integers) are handled naturally by the axis-aligned splitting rules, but unordered categorical variables must be dummy-encoded by the user. This can be suboptimal for high-cardinality factors, where native subset splits — as implemented in **bartMachine** (Kapelner and Bleich, 2016) — may be more efficient. The current AFT formulation assumes normally distributed errors; extending the framework to alternative error distributions (e.g. logistic, log-logistic) would broaden its applicability. The Gibbs sampler can be slow to mix in very high-dimensional settings; incorporating warm-starting or accelerated sampling strategies could improve computational efficiency.

ShrinkageTrees is available on CRAN and actively maintained on GitHub at `https://github.com/tijn-jacobs/ShrinkageTrees`. The methodological foundations are developed in Jacobs et al. (2025).

8 Funding

Tijn Jacobs and Stéphanie van der Pas received funding from the European Research Council (ERC) under the European Union's Horizon Europe program under Grant agreement No. 101074082. Views and opinions expressed are however those of the author(s) only and do not necessarily reflect those of the European Union or the European Research Council Executive Agency. Neither the European Union nor the granting authority can be held responsible for them.

Bibliography

Cancer Genome Atlas Research Network. Integrated genomic analyses of ovarian carcinoma. *Nature*, 474:609–615, 2011. doi: 10.1038/nature10166. [p8]

A. Caron, G. Baio, and I. Manolopoulou. Estimating individual treatment effects using non-parametric regression models: A review. *Journal of the Royal Statistical Society Series A: Statistics in Society*, 185(3): 1115–1149, 2022. doi: 10.1111/rssa.12889. [p1, 3]

C. M. Carvalho, N. G. Polson, and J. G. Scott. The horseshoe estimator for sparse signals. *Biometrika*, 97(2):465–480, 2010. doi: 10.1093/biomet/asq017. [p1, 3]

H. A. Chipman, E. I. George, and R. E. McCulloch. Bayesian CART model search. *Journal of the American Statistical Association*, 93(443):935–948, 1998. doi: 10.1080/01621459.1998.10473750. [p2]

H. A. Chipman, E. I. George, and R. E. McCulloch. BART: Bayesian additive regression trees. *The Annals of Applied Statistics*, 4(1):266–298, 2010. doi: 10.1214/09-AOAS285. [p1, 2, 3, 4]

V. Dorie. *dbarts: Discrete Bayesian Additive Regression Trees Sampler*, 2023. URL https://CRAN.R-project.org/package=dbarts. R package version 0.9-23. [p4]

V. Dorie, J. Hill, U. Shalit, M. Scott, and D. Cervone. Automated versus do-it-yourself methods for causal inference: Lessons learned from a data analysis competition. *Statistical Science*, 34(1):43–68, 2019. doi: 10.1214/18-STS667. [p1]

D. Eddelbuettel and R. François. Rcpp: Seamless R and C++ integration. *Journal of Statistical Software*, 40(8):1–18, 2011. doi: 10.18637/jss.v040.i08. [p1, 5]

P. R. Hahn, J. S. Murray, and C. M. Carvalho. Bayesian regression tree models for causal inference: Regularization, confounding, and heterogeneous effects. *Bayesian Analysis*, 15(3):965–1056, 2020. doi: 10.1214/19-BA1195. [p1, 3, 4]

F. E. Harrell, K. L. Lee, and D. B. Mark. Multivariable prognostic models: Issues in developing models, evaluating assumptions and adequacy, and measuring and reducing errors. *Statistics in Medicine*, 15 (4):361–387, 1996. doi: 10.1002/(SICI)1097-0258(19960229)15:4<361::AID-SIM168>3.0.CO;2-4. [p10]

J. He, S. Yalov, and P. R. Hahn. XBART: Accelerated Bayesian additive regression trees. In *Proceedings of the 22nd International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 89, pages 1130–1138, 2019. URL https://proceedings.mlr.press/v89/he19a.html. [p7]

J. L. Hill. Bayesian nonparametric modeling for causal inference. *Journal of Computational and Graphical Statistics*, 20(1):217–240, 2011. doi: 10.1198/jcgs.2010.08162. [p1]

T. Jacobs, W. N. van Wieringen, and S. L. van der Pas. Horseshoe forests for high-dimensional causal survival analysis. *arXiv preprint arXiv:2507.22004*, 2025. doi: 10.48550/arXiv.2507.22004. [p1, 3, 4, 13]

A. Kapelner and J. Bleich. bartMachine: Machine learning with Bayesian additive regression trees. *Journal of Statistical Software*, 70(4):1–40, 2016. doi: 10.18637/jss.v070.i04. [p4, 13]

F. Li, P. Ding, and F. Mealli. Bayesian causal inference: A critical review. *Philosophical Transactions of the Royal Society A*, 381(2247):20220153, 2023. doi: 10.1098/rsta.2022.0153. [p13]

A. R. Linero. Bayesian regression trees for high-dimensional prediction and variable selection. *Journal of the American Statistical Association*, 113(522):626–636, 2018. doi: 10.1080/01621459.2016.1264957. [p1, 3]

A. R. Linero and Y. Yang. Bayesian regression tree ensembles that adapt to smoothness and sparsity. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 80(5):1087–1110, 2018. doi: 10.1111/rssb.12293. [p4]

M. Plummer, N. Best, K. Cowles, and K. Vines. CODA: Convergence diagnosis and output analysis for MCMC. *R News*, 6(1):7–11, 2006. URL https://journal.r-project.org/archive/. [p6]

N. G. Polson and J. G. Scott. On the half-Cauchy prior for a global scale parameter. *Bayesian Analysis*, 7 (4):887–902, 2012. doi: 10.1214/12-BA730. [p3]

V. Ročková and E. Saha. On the theory of BART. *arXiv preprint arXiv:1910.00597*, 2020. [p2]

D. B. Rubin. The Bayesian bootstrap. *The Annals of Statistics*, 9(1):130–134, 1981. doi: 10.1214/aos/1176345338. [p13]

R. Sparapani, C. Spanbauer, and R. McCulloch. Nonparametric machine learning and efficient computation with Bayesian additive regression trees: The BART R package. *Journal of Statistical Software*, 97(1):1–66, 2021. doi: 10.18637/jss.v097.i01. [p1, 4, 7]

Tijn Jacobs
Vrije Universiteit Amsterdam
Department of Mathematics
Amsterdam, The Netherlands
https://tijn-jacobs.github.io/
ORCiD: 0009-0003-6188-9296
t.jacobs@vu.nl

Stéphanie van der Pas
Vrije Universiteit Amsterdam
Department of Mathematics

Amsterdam, The Netherlands
https://www.stephanievanderpas.nl/
ORCiD: 0000-0002-2448-5378
s.l.vander.pas@vu.nl

Wessel van Wieringen
Vrije Universiteit AmsterdamAmsterdam University Medical Centre
Department of Mathematics, Amsterdam, The Netherlands
Epidemiology and Data Science, Amsterdam, The Netherlands
https://www.math.vu.nl/~wvanwie/
ORCiD: 0000-0002-5100-9123
w.n.vanwieringen@vu.nl